

```
import os
import cv2
import random
import numpy as np
from glob import glob
from PIL import Image, ImageOps
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras import layers, Model
from keras import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D, UpSampling2D

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount('/content/drive', force_remount=True)

import os

# Define the path to the folder containing your ZIP file
folder_path = '/content/drive/MyDrive/ARB'

# Change the current working directory to the folder
os.chdir(folder_path)

os.listdir()

!unzip lol_dataset.zip

random.seed(5)

IMAGE_SIZE = 128
BATCH_SIZE = 4
MAX_TRAIN_IMAGES = 300

def read_image(image_path):
    image = tf.io.read_file(image_path)
    image = tf.image.decode_png(image, channels=3)
    image.set_shape([None, None, 3])
```

```
image = tf.cast(image, dtype=tf.float32) / 255.0
return image

def random_crop(low_image, enhanced_image):
    low_image_shape = tf.shape(low_image)[:2]
    low_w = tf.random.uniform(
        shape=(), maxval=low_image_shape[1] - IMAGE_SIZE + 1, dtype=tf.int32
    )
    low_h = tf.random.uniform(
        shape=(), maxval=low_image_shape[0] - IMAGE_SIZE + 1, dtype=tf.int32
    )
    enhanced_w = low_w
    enhanced_h = low_h
    low_image_cropped = low_image[
        low_h : low_h + IMAGE_SIZE, low_w : low_w + IMAGE_SIZE
    ]
    enhanced_image_cropped = enhanced_image[
        enhanced_h : enhanced_h + IMAGE_SIZE, enhanced_w : enhanced_w + IMAGE_SIZE
    ]
    return low_image_cropped, enhanced_image_cropped

def load_data(low_light_image_path, enhanced_image_path):
    low_light_image = read_image(low_light_image_path)
    enhanced_image = read_image(enhanced_image_path)
    low_light_image, enhanced_image = random_crop(low_light_image, enhanced_image)
    return low_light_image, enhanced_image

def get_dataset(low_light_images, enhanced_images):
    dataset = tf.data.Dataset.from_tensor_slices((low_light_images, enhanced_images))
    dataset = dataset.map(load_data, num_parallel_calls=tf.data.AUTOTUNE)
    dataset = dataset.batch(BATCH_SIZE, drop_remainder=True)
    return dataset

train_low_light_images = sorted(glob("./lol_dataset/our485/low/*"))[:MAX_TRAIN_IMAGES]
train_enhanced_images = sorted(glob("./lol_dataset/our485/high/*"))[:MAX_TRAIN_IMAGES]

val_low_light_images = sorted(glob("./lol_dataset/our485/low/*"))[MAX_TRAIN_IMAGES:]
val_enhanced_images = sorted(glob("./lol_dataset/our485/high/*"))[MAX_TRAIN_IMAGES:]

test_low_light_images = sorted(glob("./lol_dataset/eval15/low/*"))
test_enhanced_images = sorted(glob("./lol_dataset/eval15/high/*"))

train_dataset = get_dataset(train_low_light_images, train_enhanced_images)
```

```
val_dataset = get_dataset(val_low_light_images, val_enhanced_images)
```

```
print("Train Dataset:", train_dataset)
print("Val Dataset:", val_dataset)
```

```
Train Dataset: <_BatchDataset element_spec=(TensorSpec(shape=(4, None, None, None), dtype=tf.float32, name='image'), TensorSpec(shape=(4, None, None, None), dtype=tf.float32, name='enhanced'))>
Val Dataset: <_BatchDataset element_spec=(TensorSpec(shape=(4, None, None, None), dtype=tf.float32, name='image'), TensorSpec(shape=(4, None, None, None), dtype=tf.float32, name='enhanced'))>
```

```
# Define the residual block
```

```
def residual_block(x, filters, kernel_size):
```

```
    y = layers.Conv2D(filters, kernel_size, activation='relu', padding='same')(x)
    y = layers.BatchNormalization()(y)
    y = layers.Conv2D(filters, kernel_size, activation='relu', padding='same')(y)
    y = layers.BatchNormalization()(y)
    y = layers.Add()([x, y])
    return y
```

```
def spatial_attention_block(input_tensor):
```

```
    average_pooling = tf.reduce_max(input_tensor, axis=-1)
```

```
    average_pooling = tf.expand_dims(average_pooling, axis=-1)
```

```
    max_pooling = tf.reduce_mean(input_tensor, axis=-1)
```

```
    max_pooling = tf.expand_dims(max_pooling, axis=-1)
```

```
    concatenated = layers.concatenate([average_pooling, max_pooling])
```

```
    feature_map = layers.Conv2D(1, kernel_size=(1, 1))(concatenated)
```

```
    feature_map = tf.nn.sigmoid(feature_map)
```

```
    return input_tensor * feature_map
```

```
def channel_attention_block(input_tensor):
```

```
    channels = list(input_tensor.shape)[-1]
```

```
    average_pooling = layers.GlobalAveragePooling2D()(input_tensor)
```

```
    feature_descriptor = tf.reshape(average_pooling, shape=(-1, 1, 1, channels))
```

```
    feature_activations = layers.Conv2D(
```

```
        filters=channels // 8, kernel_size=(1, 1), activation="relu")
```

```
(feature_descriptor)
```

```
    feature_activations = layers.Conv2D(
```

```
        filters=channels, kernel_size=(1, 1), activation="sigmoid")
```

```
(feature_activations)
```

```
    return input_tensor * feature_activations
```

```
def dual_attention_unit_block(input_tensor):
```

```
    channels = list(input_tensor.shape)[-1]
```

```
    feature_map = layers.Conv2D(
```

```
    channels, kernel_size=(3, 3), padding="same", activation="relu")
)(input_tensor)
feature_map = layers.Conv2D(channels, kernel_size=(3, 3), padding="same")(
    feature_map
)
channel_attention = channel_attention_block(feature_map)
spatial_attention = spatial_attention_block(feature_map)
concatenation = layers.concatenate([channel_attention, spatial_attention])
concatenation = layers.Conv2D(channels, kernel_size=(1, 1))(concatenation)
return layers.Add()([input_tensor, concatenation])

# Add a de-blurring layer at the beginning
def deblur_layer(input_layer):
    # You can customize the parameters, such as kernel size and activation
    deblur = Conv2D(filters=64, kernel_size=3, activation='relu', padding='same')
    return deblur

def unet_model_with_multi_scale_residual_and_deblur():
    # Input Layer
    input_layer = tf.keras.layers.Input(shape=(128, 128, 3))

    # De-blurring Layer (you can adjust kernel size, stride, and filters)
    deblur = layers.Conv2D(filters=64, kernel_size=3, activation='relu', padding='same')

    # Contracting Path (Encoder) with Multi-Scale Residual Block
    conv1 = residual_block(layers.Conv2D(64, 3, activation='relu', padding='same'))
    pool1 = layers.MaxPooling2D(pool_size=(2, 2))(conv1)

    conv2 = residual_block(layers.Conv2D(128, 3, activation='relu', padding='same'))
    pool2 = layers.MaxPooling2D(pool_size=(2, 2))(conv2)

    conv3 = residual_block(layers.Conv2D(256, 3, activation='relu', padding='same'))
    pool3 = layers.MaxPooling2D(pool_size=(2, 2))(conv3)

    # Bottleneck
    conv4 = residual_block(layers.Conv2D(512, 3, activation='relu', padding='same'))

    # Expansive Path (Decoder) with Multi-Scale Residual Block
    up5 = layers.UpSampling2D(size=(2, 2))(conv4)
    up5 = layers.concatenate([up5, conv3], axis=-1)
    conv5 = residual_block(layers.Conv2D(256, 3, activation='relu', padding='same'))

    up6 = layers.UpSampling2D(size=(2, 2))(conv5)
    up6 = layers.concatenate([up6, conv2], axis=-1)
    conv6 = residual_block(layers.Conv2D(128, 3, activation='relu', padding='same'))

    up7 = layers.UpSampling2D(size=(2, 2))(conv6)
    up7 = layers.concatenate([up7, conv1], axis=-1)
```

```
conv7 = residual_block(layers.Conv2D(64, 3, activation='relu', padding='same')

# Output Layer
output_layer = layers.Conv2D(3, 1, activation='sigmoid')(conv7)

model = Model(inputs=input_layer, outputs=output_layer)

return model

# Create U-Net model with the multi-scale residual block
model = unet_model_with_multi_scale_residual_and_deblur()

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Print model summary
model.summary()

def charbonnier_loss(y_true, y_pred):
    return tf.reduce_mean(tf.sqrt(tf.square(y_true - y_pred) + tf.square(1e-3)))

def peak_signal_noise_ratio(y_true, y_pred):
    return tf.image.psnr(y_pred, y_true, max_val=255.0)

optimizer = keras.optimizers.Adam(learning_rate=0.00005)
model.compile(
    optimizer=optimizer, loss=charbonnier_loss, metrics=[peak_signal_noise_ratio]
)

history = model.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=50,
    callbacks=[
        keras.callbacks.ReduceLROnPlateau(
            monitor="val_peak_signal_noise_ratio",
            factor=0.5,
            patience=5,
            verbose=1,
            min_delta=1e-7,
            mode="max",
        )
    ],
)
```

```
plt.plot(history.history["loss"], label="train_loss")
plt.plot(history.history["val_loss"], label="val_loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Train and Validation Losses Over Epochs", fontsize=14)
plt.legend()
plt.grid()
plt.show()

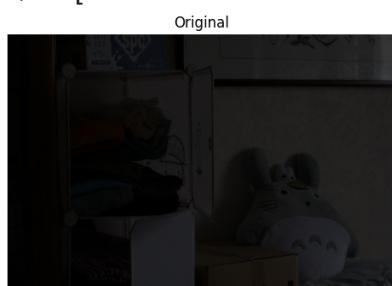
plt.plot(history.history["peak_signal_noise_ratio"], label="train_psnr")
plt.plot(history.history["val_peak_signal_noise_ratio"], label="val_psnr")
plt.xlabel("Epochs")
plt.ylabel("PSNR")
plt.title("Train and Validation PSNR Over Epochs", fontsize=14)
plt.legend()
plt.grid()
plt.show()

from tensorflow.keras.models import load_model
# Save the model if needed
model.save('/content/drive/MyDrive/ARB/image_enhancement_dunet_v250e.h5')

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3079:
      saving_api.save_model()
```

```
def plot_results(images, titles, figure_size=(12, 12)):  
    fig = plt.figure(figsize=figure_size)  
    for i in range(len(images)):  
        fig.add_subplot(1, len(images), i + 1).set_title(titles[i])  
        _ = plt.imshow(images[i])  
        plt.axis("off")  
    plt.show()  
  
def infer(original_image):  
    image = keras.utils.img_to_array(original_image)  
    image = image.astype("float32") / 255.0  
    image = np.expand_dims(image, axis=0)  
    output = model.predict(image)  
    output_image = output[0] * 255.0  
    output_image = output_image.clip(0, 255)  
    output_image = output_image.reshape(  
        (np.shape(output_image)[0], np.shape(output_image)[1], 3))  
    output_image = Image.fromarray(np.uint8(output_image))  
    original_image = Image.fromarray(np.uint8(original_image))  
    return output_image  
  
for low_light_image in random.sample(test_low_light_images, 8):  
    original_image = Image.open(low_light_image)  
    #original_image = original_image.resize((128, 128))  
    enhanced_image = infer(original_image)  
    plot_results(  
        [original_image, ImageOps.autocontrast(original_image), enhanced_image],  
        ["Original", "PIL Autocontrast", "DUO-Net Enhanced"],  
        (20, 12),  
    )
```

1/1 [=====] - 0s 40ms/step



1/1 [=====] - 0s 151ms/step



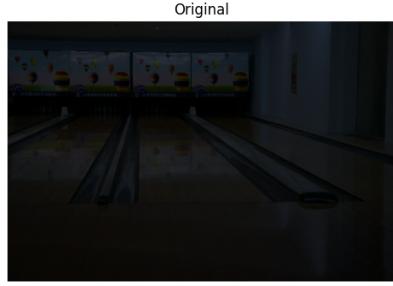
1/1 [=====] - 0s 55ms/step



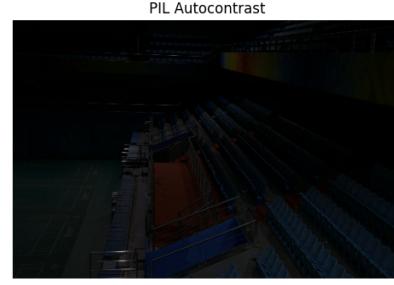
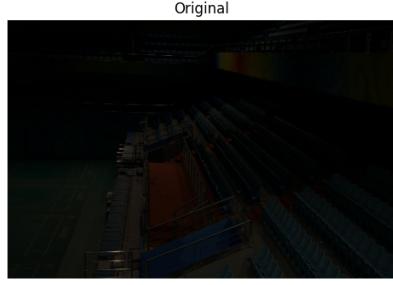
1/1 [=====] - 0s 36ms/step



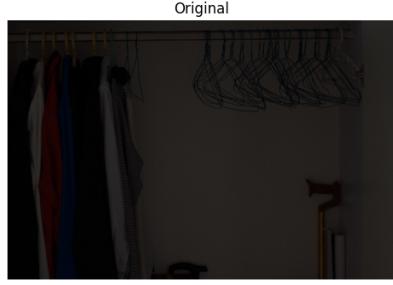
1/1 [=====] - 0s 44ms/step



1/1 [=====] - 0s 49ms/step



1/1 [=====] - 0s 60ms/step



1 / 1 [=====] - 0s 38ms/step



